


NEUROEVOLUTION

Luca Manzoni



A QUICK AND NOT-TOO-INACCURATE RECAP ON NEURAL NETWORKS


Input layer



Output layer

Hidden layer

A QUICK AND NOT-TOO-INACCURATE RECAP ON NEURAL NETWORKS



NEURAL NETWORK TRAINING

- Training is usually performed via (stochastic) gradient descend
 - Main idea: find in which direction the weights can be modified to reduce the error by differentiating the error w.r.t. the weights
 - However, only the weights are learned
 - The architecture of the network (i.e., number, size, and type of hidden layers) and the hyperparameters are selected manually
-

WHAT TO EVOLVE

- Weights (but this can also be done with backpropagation)
 - Activation functions
 - Hyperparameters (e.g., momentum, learning rate, dropout)
 - Architecture (e.g., connections, types of layer)
-

HISTORY (CLASSIC NEUROEVOLUTION)

- History: anything before the advent of large/deep networks
 - Evolution of weights and topology
 - The “unit” was the single neuron/single weight
 - Currently unfeasible due to the large number of neurons/weights/layers
-

ENCODINGS


How a neural network is represented (i.e., the genotype) is important for the kind of operations (mutation and crossover) that can be performed:

- GENITOR encoding
 - Matrix encoding
 - Node-based encoding
 - Path-based encoding
-

GENITOR ENCODING

The topology of the network is fixed

Each edge of the network is encoded as a binary string of fixed length:



Is the edge present?


Weight of the connection

101001000

The entire network can be represented as a binary string

MATRIX ENCODING

The connections can be modified...
...but not the number of neurons



The network is represented as a binary matrix


$$\begin{bmatrix} 0 & 0 & 1 & 1 & 0 \\ 0 & 0 & 0 & 1 & 0 \\ 0 & 0 & 0 & 0 & 0 \\ 0 & 0 & 0 & 0 & 1 \\ 0 & 0 & 0 & 0 & 0 \end{bmatrix}$$

Only this part is necessary
if the network is feed-forward

Does not scale even for small networks

NODE-BASED ENCODING

The main “unit” is the node, that also keeps track of the connections and the weights





Node 1	Node 2	Node 3	Node 4	Node 5
Out: [3, 4]	Out: [4]	Out: []	Out: [5]	Out: []

Both new nodes and new connections can be created

A similar encoding is used by NEAT

PATH ENCODING



The network is encoded as a collection of paths from the inputs to the outputs



We can reconstruct the network from the paths

Recombination is usually two points crossover and mutations modify a single path



THE COMPETING CONVENTIONS PROBLEM



Suppose that each hidden neuron has one of four different functionalities: A, B, C, D.

- Are the two networks different w.r.t. their represented functions?
 - What can happen when we perform crossover?
-

THE COMPETING CONVENTIONS PROBLEM



If crossover is performed without some kind of matching we can end up with two networks that are a lot worse than their parents

NEAT


NEUROEVOLUTION OF AUGMENTING TOPOLOGIES

- Think “direct encoding of a graph structure” where both the structure of the network and the weight are evolved at the same time
 - The initial population start “simple” (no hidden nodes). Evolution might add new nodes
 - Each gene has a “birth time” (called innovation number) that tracks when it was introduced. This feature is used to allow crossover
 - Features are protected with speciation. That is, at each generation only individuals in the same specie can mate
-

NEAT: THE GENOME

- **Neuron genes.** These genes encode each neuron by giving to it an ID and a type (e.g., input, output, hidden, bias, etc.).
 - **Link genes.** Each link gene contains:
 - The two endpoints of the edge
 - The weight of the edge
 - If the link is enabled or not
 - An *innovation number* used for crossover
-

NEAT: THE GENOME



Neuron Genes

ID: 1 Type: input	ID: 2 Type: input	ID: 3 Type: hidden	ID: 4 Type: hidden	ID: 5 Type: output
----------------------	----------------------	-----------------------	-----------------------	-----------------------

Link Genes

Source: 1 Destination: 5 Weight: 0.4 Enabled: Y Innovation: 3	Source: 2 Destination: 3 Weight: -0.2 Enabled: Y Innovation: 4	Source: 2 Destination: 4 Weight: 1.3 Enabled: Y Innovation: 1	Source: 3 Destination: 5 Weight: -0.9 Enabled: Y Innovation: 7	Source: 4 Destination: 5 Weight: 0.1 Enabled: Y Innovation: 6	Source: 1 Destination: 3 Weight: 1.8 Enabled: N Innovation: 5
---	--	---	--	---	---


Notice that edges might be present
in the genome but not enabled!



MUTATION(S)

- There are multiple mutations that are possible and they work on node, edges and weights:
 - Addition of a node (see next slide)
 - Add an edge with a random weight
 - Perform a slight perturbation of a weight
 - Change completely the weight of a connection
 - Enable or disable an edge
-

MUTATION: ADD NODE







- The addition of a node “splits” an edge with weight w by inserting a node in the middle
 - The outgoing edge will have weight w while the incoming one 1
-

THE INNOVATION NUMBER


- Each gene inserted by a mutation has an associated **innovation number**, a value that is always incremented and assured to be different for different mutations
 - If the **same mutation** happens multiple time in the same generation (e.g., the addition of a node between node 3 and 4), then it has the same innovation number
 - The innovation number is used during crossover and it is essential to limit the competing conventions problem
-

THE INNOVATION NUMBER



- This is the same mutation in two different individual
- The new edges $3 \rightarrow 6$ in both networks will have the same innovation number
- The same also for the new edge $6 \rightarrow 5$.


CROSSOVER





1	3	4	5	6	7
2 → 4	1 → 5	2 → 3	1 → 3	4 → 5	3 → 5

- Let us look at two individuals
- We have sorted the link genes by increasing innovation number
- We can now proceed in aligning the two genomes

1	2	4	5	7	9	11
2 → 4	2 → 5	2 → 3	1 → 3	3 → 5	3 → 4	4 → 3



CROSSOVER



- Genes in common are inherited randomly from one of the two parents
- Disjoint and excess genes are only inherited by the fittest of the parents

SPECIATION


- *A species is a group of populations with similar characteristics that are capable of successfully interbreeding with each other to produce healthy, fertile offspring, but are reproductively isolated from other species.*
 - The networks in NEAT are divided into species and crossover can only happen with individuals of the same specie
 - We need to define two aspects:
 - How new species are created
 - How species compete against each other
-

INDIRECT ENCODINGS

- Grammar-Based encoding
 - Bi-dimensional Growth encoding
 - A function giving, for each connection, its weight
 - Encoding the hyperparameters of the networks in an “aggregate” way.
E.g., “one fully connected layer with n inputs and m outputs”
-

GRAMMAR-BASED ENCODING

This is something we had already seen when working on graphs:




$$S \mapsto \begin{bmatrix} A & B \\ A & C \end{bmatrix} \quad A \mapsto \begin{bmatrix} 0 & 0 \\ 0 & 0 \end{bmatrix} \quad B \mapsto \begin{bmatrix} 1 & 0 \\ 1 & 0 \end{bmatrix} \quad C \mapsto \begin{bmatrix} 0 & 1 \\ 0 & 0 \end{bmatrix}$$

Starting from a specific *non terminal symbol*, called **axiom**, we derive, by application of production rules, the matrix representing the network

BI-DIMENSIONAL GROWTH ENCODING




L-Systems



- L-systems are a rewriting systems that can be used to create images recursively
 - At each branching point we decide
 - How many branches to create
 - The angle of each branch
 - How long is each branch with respect to the parent branch
-


BI-DIMENSIONAL GROWTH ENCODING

- We associate a L-system to each neuron
- We draw the resulting axons
- We obtain the connections of the network



CPPN


COMPOSITIONAL PATTERN-PRODUCING NETWORKS



A “network” of functions that maps a position (i, j) in a 2D plane (more dimensions are possible) to a value

This can be evolved as a “classical” network. E.g., with NEAT.
So, where's the indirect encoding?

HYPERNEAT



DENSER

DEEP EVOLUTIONARY NETWORK STRUCTURED REPRESENTATION

- A layer-based approach:
the number, type, and parameters of the layers are evolved
 - The weights are obtained via backpropagation
 - A two-levels approach:
 - the sequence of layer and a “general” type is encoded by a GA
 - The parameters of each layer are generated by grammatical evolution (GE), in particular dynamic structured GE
 - Only the best-performing networks are trained for more than a few epochs
-